



# Process control



---

# Process

---

## Process

- a program in execution
- program image + environment
  - `text(code)` / `data` / `stack` / `heap` segment
  - kernel data structure, address space, open files, ...

# Process

```
/* test_program.c */
#include <stdio.h>
int a,b;
int glob_var = 3;

void test_func(void) {
    int local_var, *buf;

    buf = (int *) malloc(10,000 * sizeof(int));
    ...
}

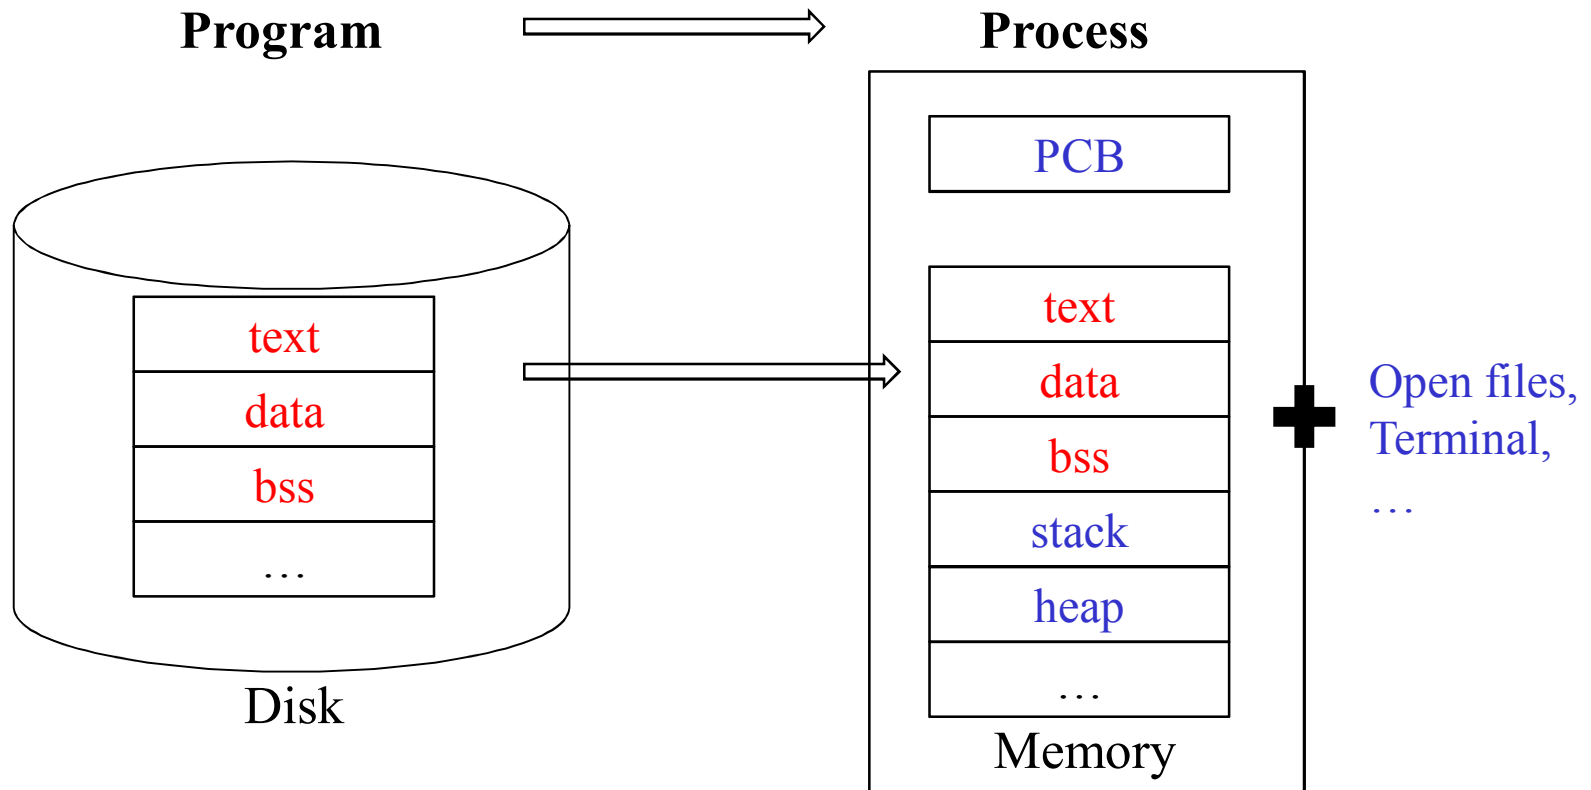
int main(int argc, char *argv[]) {
    int i = 1;
    int local_var;

    a = i + 1;
    printf("value of a = %d\n", a);
    test_func();
}
```

```
$ gcc -o test_program test_program.c
```

# Process

💻 \$ ./test\_program



\* bss: Block Started by Symbol  
PCB: Process Control Block

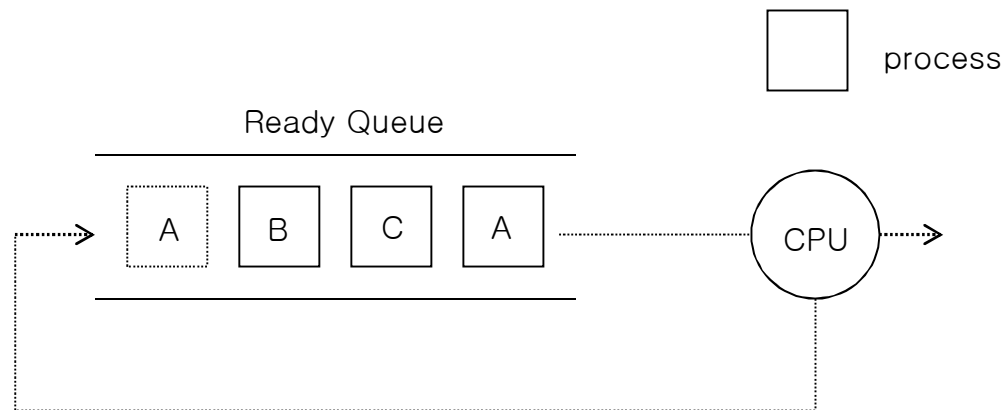
# Process

 UNIX is a multitasking system.

```
$ ps -e
PID TTY      TIME CMD
1 ?        00:00:00 init
2 ?        00:00:00 migration/0
3 ?        00:00:00 ksoftirqd/0
4 ?        00:00:00 watchdog/0
5 ?        00:00:00 migration/1
6 ?        00:00:00 ksoftirqd/1
7 ?        00:00:00 watchdog/1
8 ?        00:00:00 migration/2
...
29122 pts/9    00:00:00 bash
29151 ?        00:00:00 sshd
29170 ?        00:00:00 sshd
29171 pts/10    00:00:00 bash
30465 ?        00:00:00 httpd
$
```

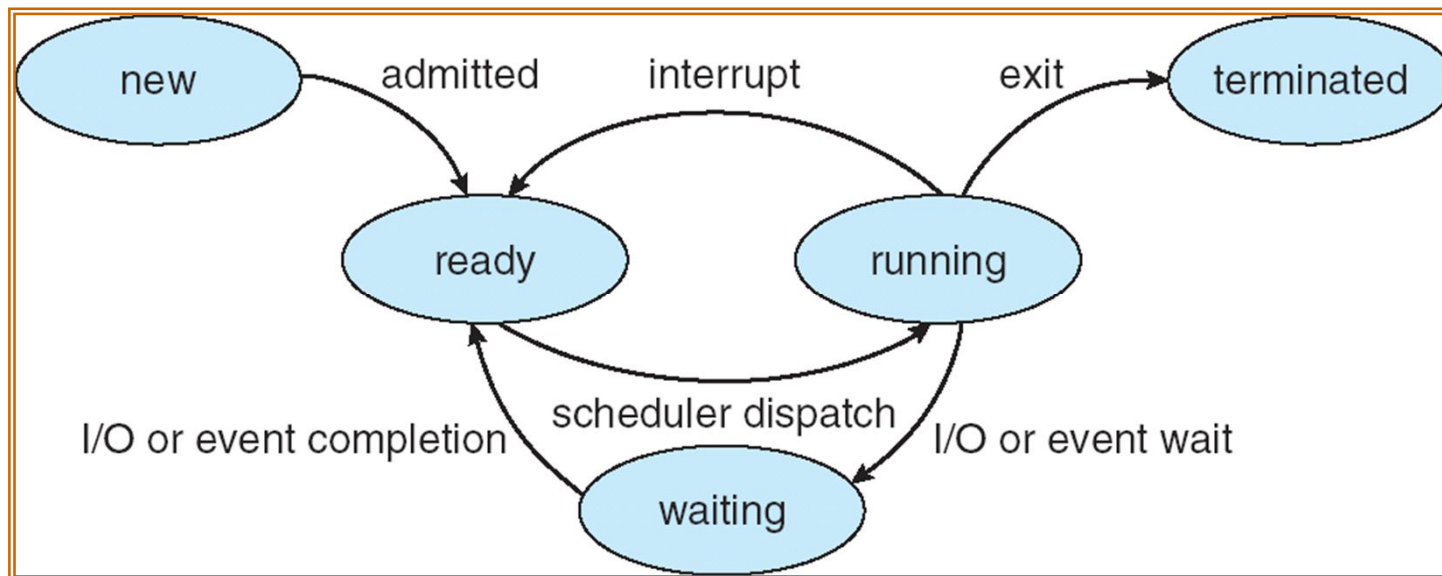
# Process

🖥️ UNIX is a time sharing system.



# Process

## State of a process



# Process identifiers

## process id

- Every process has a unique process ID.
- As processes terminate, IDs become candidates for reuse.

```
$ ps -e
PID TTY      TIME CMD
1 ?        00:00:00 init
2 ?        00:00:00 migration/0
3 ?        00:00:00 ksoftirqd/0
4 ?        00:00:00 watchdog/0
5 ?        00:00:00 migration/1
...
29122 pts/9    00:00:00 bash
29151 ?        00:00:00 sshd
29170 ?        00:00:00 sshd
29171 pts/10    00:00:00 bash
30465 ?        00:00:00 httpd
$
```



# Process identifiers

## Process 0 (scheduler process or swapper)

- Part of the kernel.
- System process.

## Process 1 (init process)

- /sbin/init
- Invoked at the end of the bootstrap procedure.
- Initialize the UNIX system with /etc/rc\* or /etc/inittab.
- Never dies.
- User process with superuser privilege.

# Process identifiers

```
#include <unistd.h>
```

```
pid_t getpid(void);
```

Returns: process ID of calling process

```
pid_t getppid(void);
```

Returns: parent process ID of calling process

```
uid_t getuid(void);
```

Returns: real user ID of calling process

```
uid_t geteuid(void);
```

Returns: effective user ID of calling process

```
gid_t getgid(void);
```

Returns: real group ID of calling process

```
gid_t getegid(void);
```

Returns: effective group ID of calling process

# fork()

```
#include <unistd.h>
```

```
pid_t fork(void);
```

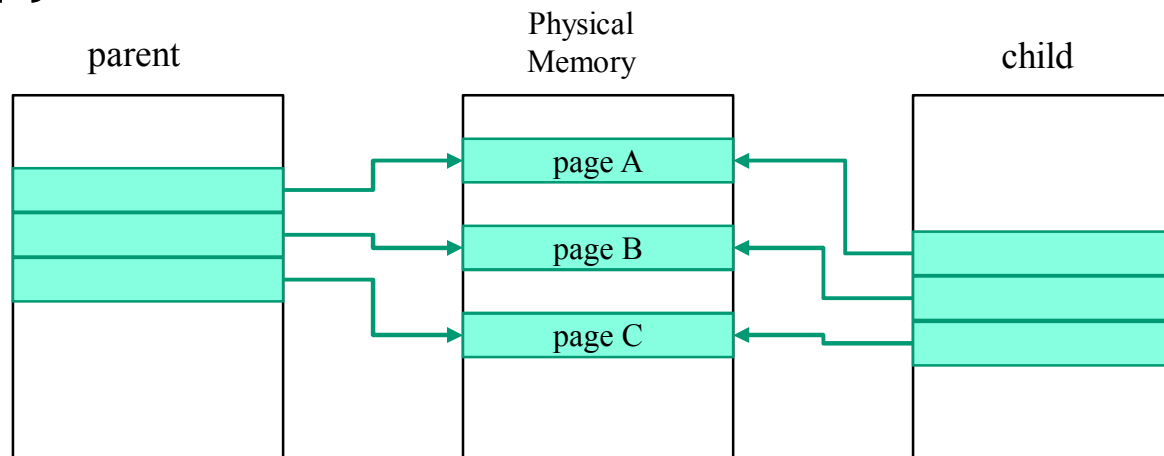
Returns: 0 in child, process ID of child in parent, -1 on error

## Create a child process.

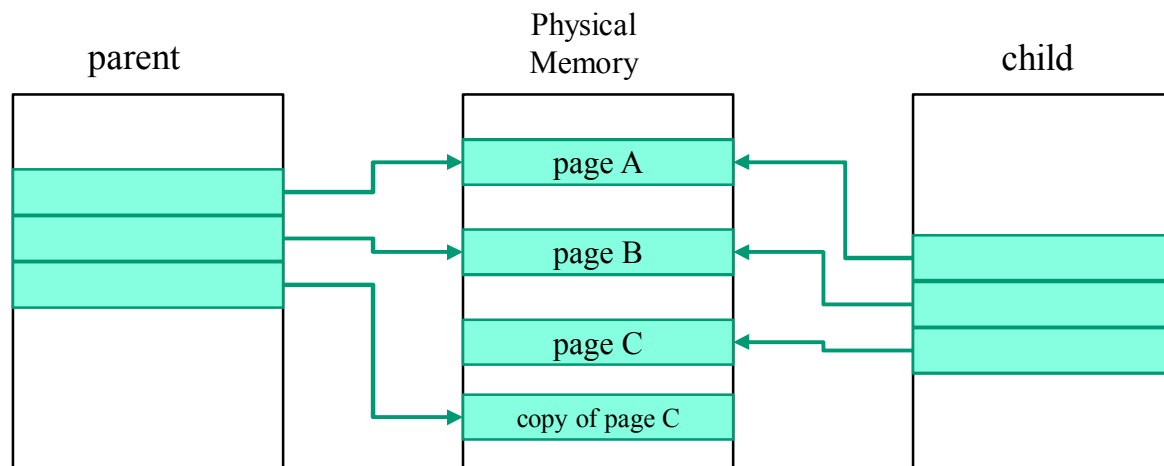
- It is called once but returns twice.
  - Return value of child: 0
  - Return value of parent: PID of child
- Both child and parent continue executing with the instruction that follows the call to fork().
- The child is a copy of parent.
  - Child gets a copy of the parent's data, heap, and stack.
  - Parent and child often share the text segment.

# fork()

## Copy on write



- After parent process write a page C.



# fork()

## Example

```
#include "apue.h"

int  glob = 6;    /* external variable in initialized data */
char buf[] = "a write to stdout\n";

int
main(void)
{
    int  var;    /* automatic variable on the stack */
    pid_t  pid;

    var = 88;
    if (write(STDOUT_FILENO, buf, sizeof(buf)-1) != sizeof(buf)-1)
        err_sys("write error");
    printf("before fork\n");    /* we don't flush stdout */
}
```

# fork()

## Example(cont.)

```
if ((pid = fork()) < 0) {
    err_sys("fork error");
} else if (pid == 0) {          /* child */
    glob++;                    /* modify variables */
    var++;
} else {
    sleep(2);                  /* parent */
}

printf("pid = %d, glob = %d, var = %d\n", getpid(), glob, var);
exit(0);
}
```

# fork()

```
#include "apue.h"
int  glob = 6;
char buf[] = "a write to stdout\n";

int main(void)
{
    int    var;
    pid_t  pid;

    var = 88;
    if (write(STDOUT_FILENO, buf, sizeof(buf)-1) != sizeof(buf)-1)
        err_sys("write error");
    printf("before fork\n"); /* we don't flush stdout */

    → if ((pid = fork()) < 0) {          /* pid is child's pid(non-0). */
        err_sys("fork error");
    } else if (pid == 0) {              /* child */
        glob++;                        /* modify variables */
        var++;
    } else {
        sleep(2);                      /* parent */
    }

    printf("pid = %d, glob = %d, var = %d\n", getpid(), glob, var);
    exit(0);
}
```

Parent process

```
#include "apue.h"
int  glob = 6;
char buf[] = "a write to stdout\n";

int main(void)
{
    int    var;
    pid_t  pid;

    var = 88;
    if (write(STDOUT_FILENO, buf, sizeof(buf)-1) != sizeof(buf)-1)
        err_sys("write error");
    printf("before fork\n"); /* we don't flush stdout */

    → if ((pid = fork()) < 0) {          /* pid is 0. */
        err_sys("fork error");
    } else if (pid == 0) {              /* child */
        glob++;                        /* modify variables */
        var++;
    } else {
        sleep(2);                      /* parent */
    }

    printf("pid = %d, glob = %d, var = %d\n", getpid(), glob, var);
    exit(0);
}
```

Child process

# fork()

## 실행

```
$ ./a.out
a write to stdout
before fork
pid = 430, glob = 7, var = 89
pid = 429, glob = 6, var = 88
$ ./a.out > temp.out
$ cat temp.out
a write to stdout
before fork
pid = 432, glob = 7, var = 89
before fork
pid = 431, glob = 6, var = 88
$
```

We never know if the child starts executing before the parent or vice versa.  
- It depends on the process scheduling algorithm.

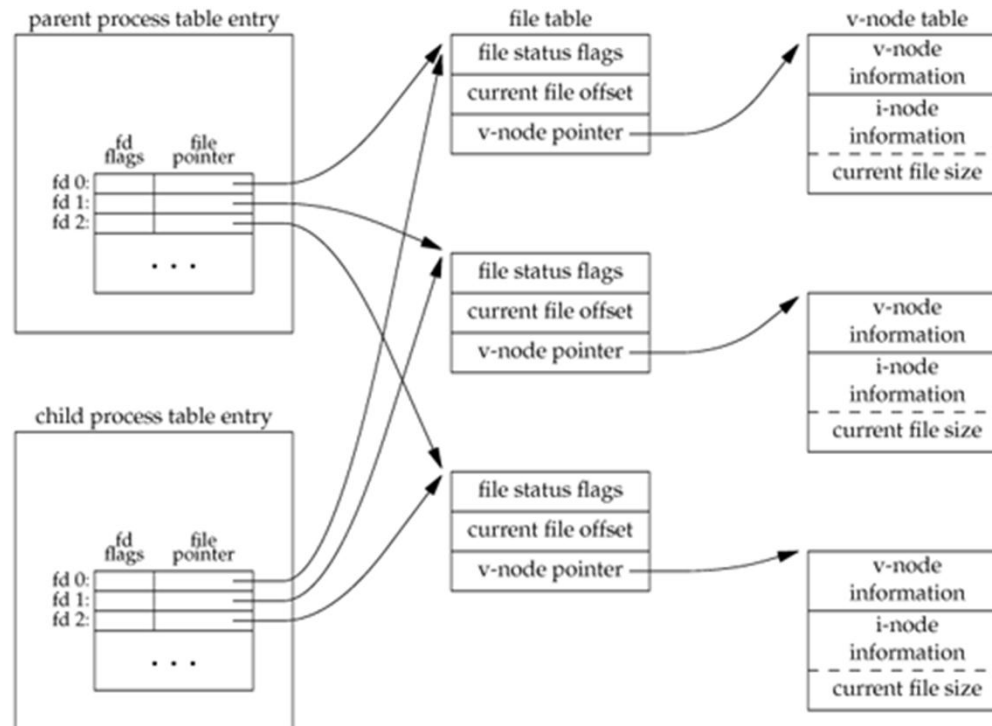


# fork()



## File sharing

- Sharing of open files between parent and child after fork()



---

# fork()

---

## Information shared by child and parent.

- real-uid(gid), effective-uid(gid)
- controlling terminal
- current working directory, root directory
- signal handlers
- environment
- resource limits

## differences between child and parent.

- the return value from fork.
- PID and PPID
- child's resource utilizations are set to 0.
- pending signals

---

# fork()

---

## Two main reasons for fork to fail

- if there are already too many processes in the system.
- if the total number of processes for this real user ID exceeds the system's limit.

## Two uses for fork

- A process wants to duplicate itself.
  - Parent and child can each execute different sections of code at the same time.
  - Common for [network servers](#).
- A process wants to execute a different program.
  - Common for shells.
  - Child does [exec\(\)](#) right after it returns from [fork\(\)](#).

# vfork()

```
#include <unistd.h>
```

```
pid_t vfork(void);
```

Returns: 0 in child, process ID of child in parent, -1 on error

- same calling sequence and same return values as fork().
- intended to create a new process when the purpose of the new process is to exec() a new program.
  - Does not copy the address space of parent into the child.
  - The child calls exec() or exit() right after the vfork().
  - The child runs in the address space of the parent.
  - Provides an efficiency.
- vfork() guarantees that the child runs first.

# vfork()

## Example

```
#include "apue.h"
int  glob = 6;    /* external variable in initialized data */

int
main(void)
{
    int  var;      /* automatic variable on the stack */
    pid_t pid;

    var = 88;
    printf("before vfork\n"); /* we don't flush stdio */
    if ((pid = vfork()) < 0) {
        err_sys("vfork error");
    } else if (pid == 0) {    /* child */
        glob++;              /* modify parent's variables */
        var++;
        exit(0);             /* child terminates */
    }
}
```

# vfork()

## Example(cont.)

```
/*  
 * Parent continues here.  
 */  
printf("pid = %d, glob = %d, var = %d\n", getpid(), glob, var);  
exit(0);  
}
```


## 실행

```
$ ./a.out  
before vfork  
pid = 29039, glob = 7, var = 89  
$
```

# exit()

```
#include <stdlib.h>
void exit(int status);
void _Exit(int status);

#include <unistd.h>
void _exit(int status);
```

 causes **normal program termination** and the value of status is returned to the parent.

 **\_exit** and **\_Exit**

- return to the kernel immediately.

 **exit**

- performs cleanup processing and then returns to the kernel.
  - all open streams are flushed and closed.

# Process termination

## Two types of process termination

- **normal** termination
  - return from main
  - calling exit
  - calling \_exit or \_Exit
- **abnormal** termination
  - calling abort
  - Receipt of a signal

## Regardless of how a process terminate, the same code is eventually executed.

- Close open descriptors, release the memory, ...



# Process termination

## Termination status

- Terminating process notify its parent how it terminated.
- normal termination
  - pass an exit status as argument to `exit()` or `_exit()`.
  - return value from `main()`
  - exit status is converted into a termination status.
- abnormal termination
  - Kernel generates a termination status to indicate the reason for the abnormal termination.
- The parent of the terminated process can obtain the termination status from **`wait()` or `waitpid()`**.

# Process termination

## zombie state

- Kernel has to keep some information
  - PID, termination status, and CPU usage time
  - This information is available when parent calls `wait()`.
- the process that has terminated, but whose parent has not yet waited for it, is called a **zombie**.

## If the parent terminates before the child?

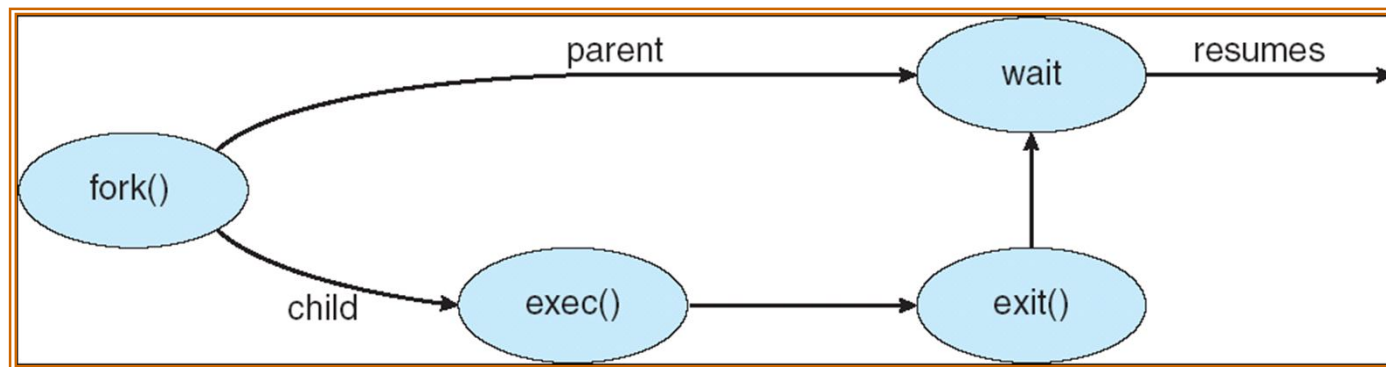
- `init` process becomes the parent of the orphaned process.

## If an orphaned process is terminated?

- `init` calls `wait()` to fetch the termination status.

# wait() and waitpid()

- ❏ A process that calls wait() or waitpid()
- Block, if all of its children are still running.
  - Return immediately if a child has terminated.
  - Return immediately with an error, if it doesn't have any child processes.



# wait() and waitpid()

```
#include <sys/wait.h>
```

```
pid_t wait(int *statloc);
```

**Both return: process ID if OK, 0 (see later), or -1 on error**

## wait()

- If a child has already terminated and is a zombie, wait returns immediately with that child's status.
- Otherwise, it blocks the caller until a child terminates.
- If the caller blocks and has multiple children, wait returns when one terminates.

## statloc argument

- Store termination status in the location pointed to by *statloc*.

# wait() and waitpid()

## Macro to examine the termination status

- WIFEXITED(status)
  - 정상적으로 종료한 경우에 참 값을 리턴
- WEXITEDSTATUS(status)
  - exit 함수의 인자에서 하위 8 비트 값을 리턴
- WIFSIGNALED(status)
  - 시그널에 의해 종료한 경우에 참 값을 리턴
- WTERMSIG(status)
  - 시그널 번호를 리턴
- WIFSTOPPED(status)
  - 실행이 일시 중단된 경우에 참 값을 리턴
- WSTOPSIG(status)
  - 실행을 일시 중단시킨 시그널 번호를 리턴

# wait() and waitpid()

## Example

```
#include <sys/types.h>
#include <sys/wait.h>


void pr_exit(int status)
{
    if (WIFEXITED(status))
        printf("normal termination, exit status = %d\n",
               WEXITSTATUS(status));
    else if (WIFSIGNALED(status))
        printf("abnormal termination, signal number = %d\n",
               WTERMSIG(status));
    else if (WIFSTOPPED(status))
        printf("child stopped, signal number = %d\n",
               WSTOPSIG(status));
}
```

# wait() and waitpid()

## Example(cont.)

```
#include <sys/types.h>
#include <sys/wait.h>

int main(void)
{
    pid_t  pid;
    int    status;

    if ((pid = fork()) < 0) 
        printf("fork error");
    else if (pid == 0)          /* child */
        exit(7);

    if (wait(&status) != pid)  /* wait for child */
        printf("wait error");
    pr_exit(status);           /* and print its status */
}
```

# wait() and waitpid()

## Example(cont.)

```
if ( (pid = fork()) < 0)
    printf("fork error");
else if (pid == 0)           /* child */
    abort();                 /* generates SIGABRT */

if (wait(&status) != pid)    /* wait for child */
    printf("wait error");
pr_exit(status);            /* and print its status */

if ( (pid = fork()) < 0)
    printf("fork error");
else if (pid == 0)           /* child */
    status /= 0;             /* divide by 0 generates SIGFPE */

if (wait(&status) != pid)    /* wait for child */
    printf("wait error");
pr_exit(status);            /* and print its status */
}
```



# wait() and waitpid()

## 실행

```
$ ./a.out
normal termination, exit status = 7
abnormal termination, signal number = 6
abnormal termination, signal number = 8
$
```

Each signal number is defined in <signal.h>  
SIGABRT: 6  
SIGFPE: 8

# wait() and waitpid()

```
#include <sys/wait.h>
```

```
pid_t waitpid(pid_t pid, int *statloc, int options);
```

Both return: process ID if OK, 0 (see later), or -1 on error

## Why waitpid() is required?

- If we have more than one child, `wait()` returns on termination of any of the children.
- What if we want to wait for a specific process to terminate?

## waitpid()

- Waits for a specific process with *pid* argument.
- Provides some controls with *options* argument.

# wait() and waitpid()

## pid argument

- pid == -1
  - Waits for any child process.
  - In this case, waitpid() is equivalent to wait().
- pid > 0
  - Waits for the child whose process ID equals pid.
- pid == 0
  - Waits for any child whose process group ID equals that of the calling process.
- pid < -1
  - Waits for any child whose process group ID equals the absolute value of pid.

---

## wait() and waitpid()

---

### options argument

- WNOHANG
  - Not block(return immediately) if a child specified by pid is not terminated.
- WUNTRACED
  - Not block if a child specified by pid has stopped.

# wait() and waitpid()

## Example

```
#include "apue.h"
#include <sys/wait.h>

int main(void)
{
    pid_t  pid;

    if ((pid = fork()) < 0) {
        err_sys("fork error");
    } else if (pid == 0) {          /* first child */
        if ((pid = fork()) < 0)
            err_sys("fork error");
        else if (pid > 0)
            exit(0);               /* parent from second fork == first child */
        /* We're the second child; our parent becomes init. */
        sleep(2);
        printf("second child, parent pid = %d\n", getppid());
        exit(0);
    }
}
```

# wait() and waitpid()

## Example

```
if (waitpid(pid, NULL, 0) != pid) /* wait for first child */
    err_sys("waitpid error");

/*
 * We're the parent (the original process); we continue executing,
 * knowing that we're not the parent of the second child.
 */
exit(0);
}
```

## 실행

```
$ ./a.out
$ second child, parent pid = 1
```

# exec()

```
#include <unistd.h>
```

```
int execl(const char *pathname, const char *arg0, ... /* (char *)0 */);
```

```
int execv(const char *pathname, char *const argv []);
```

```
int execl(const char *pathname, const char *arg0, ... /* (char *)0,  
            char *const envp[] */);
```

```
int execve(const char *pathname, char *const argv[], char *const envp []);
```

```
int execlp(const char *filename, const char *arg0, ... /* (char *)0 */);
```

```
int execvp(const char *filename, char *const argv []);
```

All six return: -1 on error, no return on success

## Execute a program.

- When a process calls one of the exec functions, the process is completely replaced by the new program.
- New program starts executing at its main() function.
- exec() merely replaces the current process(text, data) with a brand new program from disk.

# exec()

## pathname(filename) argument

- pathname(filename) of a file to be executed.
- execl, execv, execl, execve take a pathname argument.
- execl<sub>p</sub>, execv<sub>p</sub> take a filename argument.
  - If filename contains a '/', it is taken as a pathname.
  - Otherwise, the executable file is searched for in the directories specified by PATH.
    - E.g. PATH=/bin:/usr/bin:/usr/local/bin/::



# exec()

## Argument list

- `arg0, arg1, ..., argn` in the `execl`, `execlp`, and `execle`.
- The list of arguments is terminated by a NULL pointer.
  - `execl("/bin/ls", "ls", "-al", 0);`

## Argument vector

- `*argv[]` in `execv`, `execve` and `execvp`.
- The array of pointers is terminated by a NULL pointer.
  - `char *argv[3] = {"ls", "-al", 0};`
  - `execv("/bin/ls", argv);`

# exec()

## Environment variable

- \*envp[] in `execv` and `execle`
- A pointer to an array of pointers to the environment strings.
- The other functions use the *environ* variable.
  - `char *env[2] = {"TERM=vt100", 0};`
  - `execle("/bin/ls", "ls", 0, env);`

# exec()


## Differences among the six exec functions

function	pathname	filename	arg list	argv[]	environ	envp[]
execl	0		0		0	
execlp		0	0		0	
execle	0		0			0
execv	0			0	0	
execvp		0		0	0	
execve	0			0		0
(letter in name)		p	l	v		e

# exec()

❏ `execve` is a system call.

- `execl`, `execv`, `execle`, `execlp`, `execvp` are library functions.
- Relationship of the six exec functions.



# exec()

## Example

```
#include "apue.h"
#include <sys/wait.h>

char  *env_init[] = { "USER=unknown", "PATH=/tmp", NULL };

int
main(void)
{
    pid_t  pid;

    if ((pid = fork()) < 0) {
        err_sys("fork error");
    } else if (pid == 0) { /* specify pathname, specify environment */
        if (execle("/home/sar/bin/echoall", "echoall", "myarg1",
                  "MY ARG2", (char *)0, env_init) < 0)
            err_sys("execle error");
    }
}
```

# exec()

## Example(cont.)

```
if (waitpid(pid, NULL, 0) < 0)
    err_sys("wait error");

if ((pid = fork()) < 0) {
    err_sys("fork error");
} else if (pid == 0) { /* specify filename, inherit environment */
    if (execlp("echoall", "echoall", "only 1 arg", (char *)0) < 0)
        err_sys("execlp error");
}

exit(0);
}
```

# exec()

## echoall

```
#include "apue.h"

int
main(int argc, char *argv[])
{
    int    i;
    char    **ptr;
    extern char **environ;

    for (i = 0; i < argc; i++)                /* echo all command-line args */
        printf("argv[%d]: %s\n", i, argv[i]);

    for (ptr = environ; *ptr != 0; ptr++)      /* and all env strings */
        printf("%s\n", *ptr);

    exit(0);
}
```

# exec()

## 실행

```
$ ./a.out  
argv[0]: echoall  
argv[1]: myarg1  
argv[2]: MY ARG2  
USER=unknown  
PATH=/tmp  
$ argv[0]: echoall  
argv[1]: only 1 arg  
USER=sar  
LOGNAME=sar  
SHELL=/bin/bash  
...  
HOME=/home/sar
```