

---

# Network IPC: Sockets

---


# IPC (Inter Process Communication)

🖥️ pipes, FIFOs, message queues, semaphores, and shared memory

- allow processes running on the **same machine** to communicate with one another.

🖥️ socket

- allows processes running on **different machines** to communicate with one another.



---

# Socket

---

## Socket

- an abstraction of a communication endpoint.

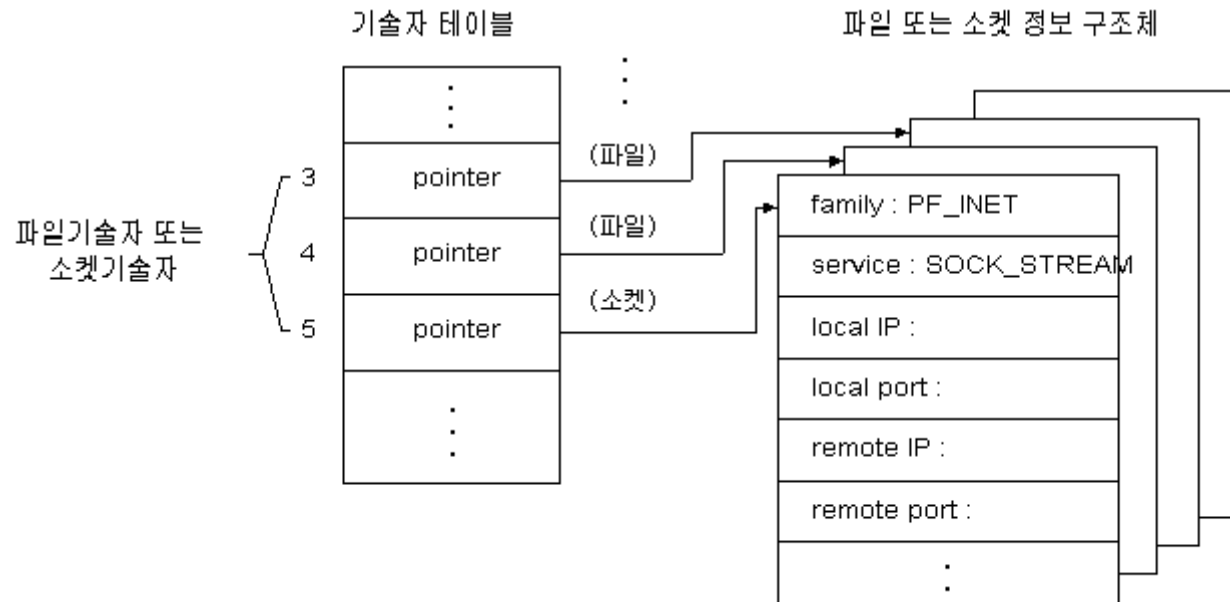
## Socket descriptor

- Just as applications would use file descriptors to access a file, applications use **socket descriptors** to access sockets.
- Socket descriptors are implemented **as file descriptors**.
  - many of the functions that deal with file descriptors (**read** and **write**) will work with a socket descriptor.
  - But, **lseek** doesn't work with sockets, since sockets don't support the concept of a file offset.

# Socket

❏ Socket descriptor and file descriptor **shares a table**.

- 3 and 4 are file descriptors.
- 5 is a socket descriptor.



- descriptor table은 per process structure이다. 따라서, 다른 process간에는 같은 값의 descriptor를 가질 수 있다.

# socket()

```
#include <sys/socket.h>
```

```
int socket(int domain, int type, int protocol);
```

Returns: file (socket) descriptor if OK, -1 on error

## Create a socket.

- *domain* argument determines the nature of the communication, including the [address format](#).
  - **AF\_INET**: IPv4 Internet domain
    - Usually, this is used.
    - **PF\_INET** is also used.
  - AF\_INET6: IPv6 Internet domain
  - AF\_UNIX: UNIX domain
  - AF\_UNSPEC: unspecified

---

# socket()

---

- *type* argument determines the type of the socket, which further determines the **communication characteristics**.
  - **SOCK\_DGRAM**: connectionless, unreliable
  - **SOCK\_STREAM**: connection-oriented, reliable
- *protocol* argument selects the default protocol for the given domain and socket type.
  - TCP (Transmission Control Protocol):
    - SOCK\_STREAM in AF\_INET domain
  - UDP (User Datagram Protocol):
    - SOCK\_DGRAM in AF\_INET domain
  - **Usually, zero is used.**

# socket()

## example

```
#include <stdio.h>
#include <sys/types.h>
#include <sys/stat.h>
#include <fcntl.h>
#include <sys/socket.h>

int main()
{
    int fd1, fd2, sd1, sd2;

    fd1 = open("/etc/passwd", O_RDONLY);
    printf("/etc/passwd's file descriptor = %d\n", fd1);

    sd1 = socket(PF_INET, SOCK_STREAM, 0);
    printf("stream socket descriptor = %d\n", sd1);
```

# socket()

## example(cont.)

```
sd2 = socket(PF_INET, SOCK_DGRAM, 0);
printf("datagram socket descriptor = %d\n", sd2);

fd2 = open("/etc/hosts", O_RDONLY);
printf("/etc/hosts's file descriptor = %d\n", fd2);

close(fd2) ;
close(fd1) ;
close(sd2) ;
close(sd1) ;

}
```



# socket()

## 실행

```
$gcc -o open_socket open_socket.c -lsocket
$./open_socket
/etc/passwd's file descriptor = 3
stream socket descriptor = 4
datagram socket descriptor = 5
/etc/hosts's file descriptor = 6
$
```

# shutdown()

```
#include <sys/socket.h>
```

```
int shutdown (int sockfd, int how);
```

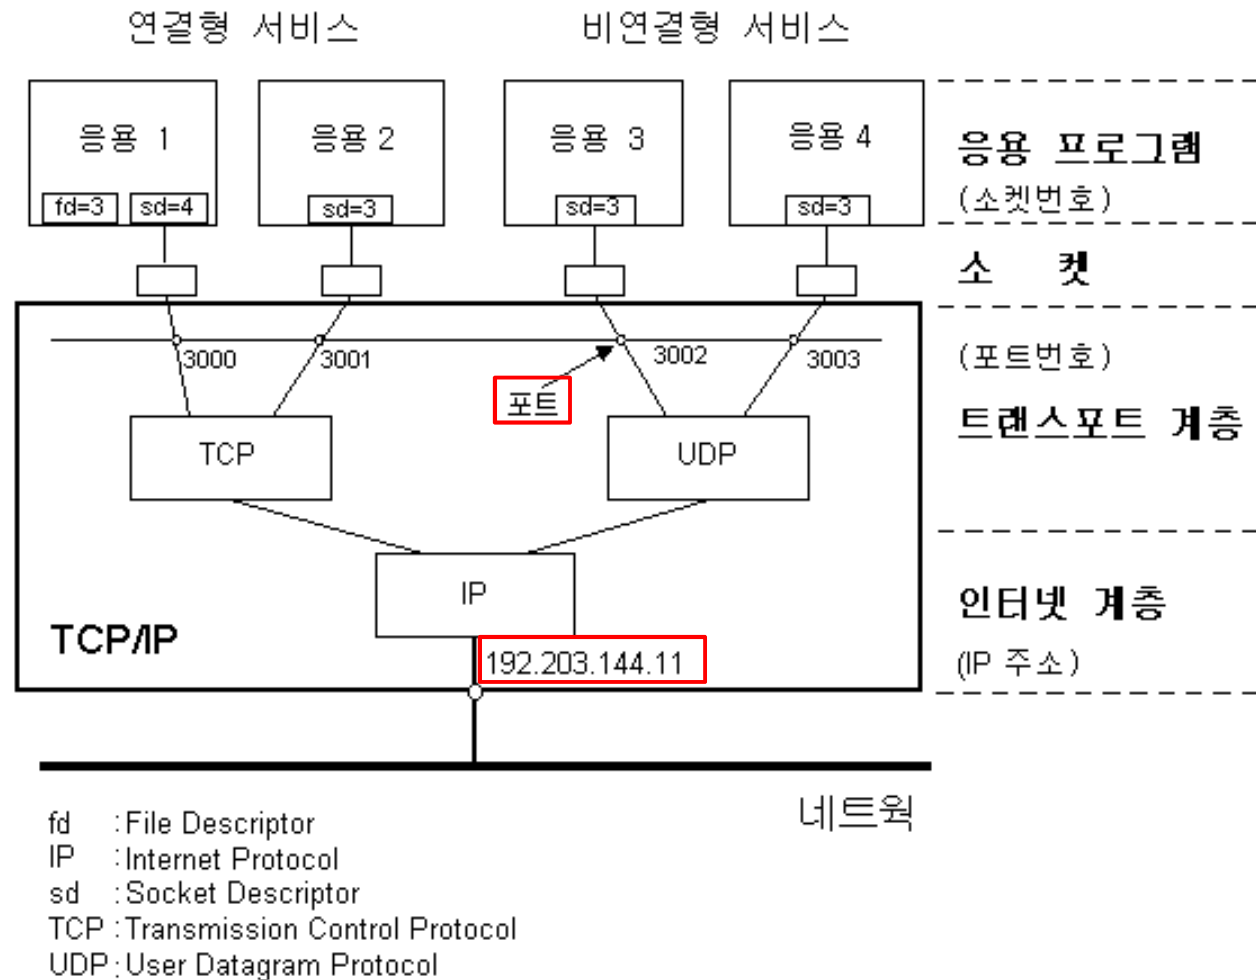
Returns: 0 if OK, -1 on error

 disable I/O on a socket.

- Is similar to [close\(\)](#).
- *how* argument
  - SHUT\_RD: reading from socket is disabled.
  - SHUT\_WR: transmitting data to socket is disabled.
  - SHUT\_RDWR: both data transmission and reception are disabled.

# Socket address

## Overall structure including socket



# Socket address

## Generic socket address structure

```
#include <sys/socket.h>

struct sockaddr {
    sa_family_t sa_family;    /* address family */
    char        sa_data[];    /* variable-length address */
};
```

- socket을 이용할 통신 객체(client or server)의 주소를 표현하기 위해서는 address family, IP address, port number가 지정되어야 하며, 이 주소 정보를 **socket address**라고 부른다.
- 이 sockaddr 구조체에 IP address, port number 등을 직접 쓰거나 읽기가 불편하므로, sockaddr 구조체를 사용하는 대신 4 바이트의 IP address와 2 바이트의 port 번호를 구분하여 지정할 수 있는 인터넷 전용 소켓주소 구조체인 sockaddr\_in을 주로 사용한다.

# Socket address

## Internet protocol address structure

```
#include <netinet/in.h>

struct in_addr {
    in_addr_t    s_addr;          /* IPv4 address */
};

struct sockaddr_in {
    sa_family_t  sin_family;      /* address family */
    in_port_t    sin_port;        /* port number (2bytes) */
    struct in_addr sin_addr;      /* IPv4 address (4bytes) */
    char sin_zero[8];            /* not used */
};
```

- sin\_family
  - AF\_INET, AF\_UNIX, ...

# Socket address

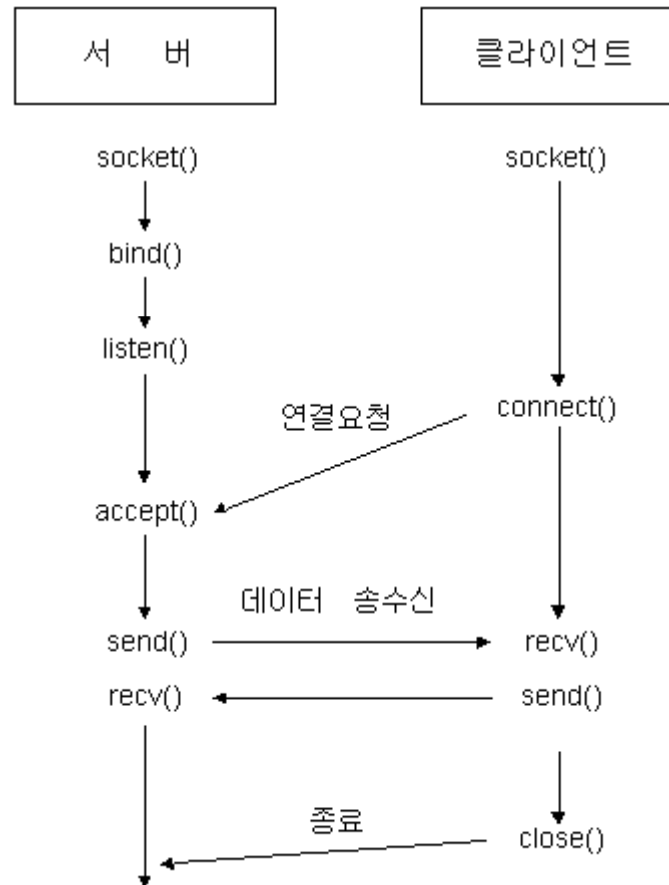
## Socket address structure

- sockaddr\_in은 sockaddr 구조체의 데이터를 internet protocol에서 사용하기에 적합하도록 수정한 것이다.

	sa_family		sa_data	
sockaddr	체 계	데 이 터		
	2바이트	2바이트	4바이트	8바이트
sockaddr_in	체 계	포 트	IP 주 소	사 용 되 지 않 음
	sin_family	sin_port	sin_addr	sin_zero

# Socket programming

## socket programming's basic procedure



# bind()

```
#include <sys/socket.h>
```

```
int bind(int sockfd, const struct sockaddr *addr, socklen_t len);
```

Returns: 0 if OK, -1 on error

 associate an address with a socket.

- For a **server**, we need to associate a **well-known address** with the **server's socket** on which client requests will arrive.
- For client, we can let the system choose a default address.
  - ➔ bind() is not used in client side.
- *len* is the size of socket address.



# connect()

```
#include <sys/socket.h>
```

```
int connect(int sockfd, const struct sockaddr *addr, socklen_t len);  
Returns: 0 if OK, -1 on error
```

 **Client** requests a connection to server.

- Before exchanging data, we need to create a connection between the socket of client and server.
- *addr* is the address of the server.

# listen()

```
#include <sys/socket.h>
```

```
int listen(int sockfd, int backlog);
```

Returns: 0 if OK, -1 on error

- ❏ A **server** announces that it is willing to accept connect requests.
  - *backlog* specifies how many connection requests can be queued. (queue size)

# accept()

```
#include <sys/socket.h>
```

```
int accept(int sockfd, struct sockaddr *addr, socklen_t *len);
```

Returns: file (socket) descriptor if OK, -1 on error

 **Server** accept a connect request from clients.

- *addr* is the **socket address of client**.
- The file descriptor returned is a socket descriptor that is connected to the client that called connect().

# send()

```
#include <sys/socket.h>
```

```
ssize_t send(int sockfd, const void *buf, size_t nbytes, int flags);
```

Returns: number of bytes sent if OK, -1 on error

 Send data to other process.

- is similar to [write](#), but
- allows us to specify flags to change how the data we want to transmit is treated
  - MSG\_DONTROUTE, MSG\_DONTWAIT, MSG\_EOR, MSG\_OOB

# receive()

```
#include <sys/socket.h>
```

```
ssize_t recv(int sockfd, void *buf, size_t nbytes, int flags);
```

**Returns:** length of message in bytes,

**0** if no messages are available and peer has done an orderly shutdown,  
**or -1** on error

 Receive data from other process.

- is similar to [read](#), but
- allows us to specify some options to control how we receive the data.
  - MSG\_OOB, MSG\_PEEK, MSG\_TRUNC, MSG\_WAITALL

# Byte ordering

## Two types of byte order

- Little endian
  - the least significant byte contains the lowest byte address
  - Intel, DEC
- Big endian
  - the least significant byte contains the highest byte address
  - IBM, Motorola
- Eg. 0xC3E2

주소 :	n	n+1
데이터 :	E2	C3

(a) 80x86 계열

n	n+1
C3	E2

(b) MC68000계열

# Byte ordering

## Byte ordering

- Host byte order
  - 컴퓨터가 memory에 byte를 저장하는 순서
- Network byte order
  - network에서 byte 단위로 data가 전달되는 순서
  - 0xC3E2의 경우 C3, E2의 순서로 전달 (like big endian)
- 즉, intel의 80x86계열의 CPU가 사용하는 host byte order는 network byte order와 다르다. 따라서 80x86계열의 컴퓨터에서 network를 통하여 전송한 데이터를 68000계열의 컴퓨터가 수신하면 byte 순서가 바뀌게 된다.

# Byte ordering

## Byte ordering


- 이러한 문제를 해결하기 위하여 컴퓨터 내부에서 만들어진 **host byte order** 데이터를 **network**로 전송하기 전에 **htons()** 함수를 사용하여 모두 **network byte order**로 바꾸어야 한다.
- 반대로 **network**에서 수신한 데이터는 **ntohs()** 함수를 사용하여 자신에게 맞는 **host byte order**로 바꾸어야 한다.
- 즉, **network byte order**를 지켜 데이터를 전송함으로써 수신한 데이터가 어떤 종류의 컴퓨터에서 만들어진 것인지 알 필요가 없도록 하는 것이다.
- **Motolora** 계열의 **CPU**에서는 **host byte order**와 **network byte order**가 같은데, 이러한 호스트에서의 **htons()**와 **ntohs()** 함수는 아무 일도 하지 않는다.



# Byte ordering

```
#include <arpa/inet.h>
```

```
uint32_t htonl(uint32_t hostint32);  
uint16_t htons(uint16_t hostint16);  
uint32_t ntohl(uint32_t netint32);  
uint16_t ntohs(uint16_t netint16);
```

 convert between the host byte order and the network byte order.

- htonl/htons (4byte/2byte)
  - Host byte order → network byte order
- ntohl/ntohs (4byte/2byte)
  - Network byte order → host byte order

# Address conversion

## IP address conversion

- 32비트의 IP address를 편의에 따라 sp.kwangwoon.ac.kr과 같은 domain name, 그리고 192.203.144.11과 같은 dotted decimal 표시법 등으로 바꾸어 사용하고 있다.
- 한편 IP packet을 network로 실제로 전송할 때에는 32비트의 IP address가 필요하다.
- 따라서, 이들 주소 표현법을 자유롭게 변환할 수 있는 함수가 필요하다.

# Address conversion

## IP address conversion

- 예: 아래 그림에서 dotted decimal로 표현된 192. 203.144.11을 32 비트의 IP address로 변환하려면 `inet_addr()` 시스템 콜을 사용하고 IP address를 다시 dotted decimal로 변환하려면 `inet_ntoa()`를 사용한다.

sp.kwangwoon.ac.kr    

11000000	11001011	10010000	00001011
----------	----------	----------	----------

 : 192.203.144.11

도메인 네임                      :                      IP 주소 (binary)                      : dotted decimal

gethostbyname()  
→  
←  
gethostbyaddr()

inet\_addr()  
←  
→  
inet\_ntoa()

# Address conversion

## example

```
struct sockaddr_in server_addr;  
server_addr.sin_addr.s_addr = inet_addr ("192.203.144.11");  
printf("%x\n", server_addr.sin_addr.s_addr);           /* hexa 4바이트 출력 */  
printf("%s\n", inet_ntoa(server_addr.sin_addr));       /* dotted decimal 출력 */
```

# Server example

## Example

```
#include <sys/types.h>
#include <sys/socket.h>
#include <unistd.h>
#include <arpa/inet.h>
#include <stdio.h>
#include <stdlib.h>
#include <string.h>

#define PORT 5555

int main(void)
{
    char buf[256];
    struct sockaddr_in server, client;
    int sd, cd, clientlen = sizeof(client);
```

# Server example

## Example(cont.)

```
if ((sd = socket(AF_INET, SOCK_STREAM, 0)) == -1) {
    perror("socket");
    exit(1);
}
memset((char *)&server, '\0', sizeof(server));
server.sin_family      = AF_INET;
server.sin_addr.s_addr = inet_addr("127.0.0.1");
server.sin_port        = htons(PORT);

if(bind(sd, (struct sockaddr*)&server, sizeof(server))) {
    perror("bind");
    exit(1);
}
if(listen(sd, 5)) {
    perror("listen");
    exit(1);
}
```

# Server example

## Example(cont.)

```
    if ((cd = accept(sd, (struct sockaddr*)&client, &clientlen)) == -1) {  
        perror("accept");  
        exit(1);  
    }  
  
    sprintf(buf, "Your IP address is %s", inet_ntoa(client.sin_addr));  
    if(send(cd, buf, strlen(buf) + 1, 0) == -1) {  
        perror("send");  
        exit(1);  
    }  
    close(cd);  
    close(sd);  
  
    return 0;  
}
```

# Client example

## Example

```
#include <sys/types.h>
#include <sys/socket.h>
#include <unistd.h>
#include <arpa/inet.h>
#include <stdio.h>
#include <stdlib.h>
#include <string.h>

#define PORT 5555

int main(void)
{
    int sd;
    char buf[256];
    struct sockaddr_in server;
```



# Client example

## Example(cont.)

```
if ((sd = socket(AF_INET, SOCK_STREAM, 0)) == -1) {  
    perror("socket");  
    exit(1);  
}  
  
memset((char *)&server, '\0', sizeof(server));  
server.sin_family      = AF_INET;  
server.sin_addr.s_addr = inet_addr("127.0.0.1");  
server.sin_port        = htons(PORT);  
  
if(connect(sd, (struct sockaddr*)&server, sizeof(server))) {  
    perror("connect");  
    exit(1);  
}
```

# Client example

## Example(cont.)

```
    if(recv(sd, buf, sizeof(buf), 0) == -1) {  
        perror("recv");  
        exit(1);  
    }  
    close(sd);  
    printf("From Server: %s\n", buf);  
  
    return 0;  
}
```

# Server & client example

## Server 실행

```
$ ./test_server
```

(1) Server부터 실행


## Client 실행

```
$ ./test_client
```

From Server: Your IP address is 127.0.0.1

```
$
```

(2) Client 실행



```
$ ./test_server
```

```
$
```

(3) Server 종료

---

# Reference

---

## More about network programming

- UNIX Network Programming, Volume 2, Second Edition: Interprocess Communications, Prentice Hall, 1999.
- UNIX Network Programming, Volume 1, Second Edition: Networking APIs: Sockets and XTI, Prentice Hall, 1998

