



Threads



Threads concepts


Thread

- An independent and schedulable execution unit.
- A process can be divided into **two or more running threads**.
- A single thread of control(= a UNIX process)
 - Each process is doing **only one thing at a time**.
- Multiple threads of control in a process.
 - The process can do **more than one thing at a time**.
- Multithreading is possible on even uni-processor
 - by **time-division multiplexing**.

Threads concepts

📖 A typical example

- Apache web server



Threads concepts

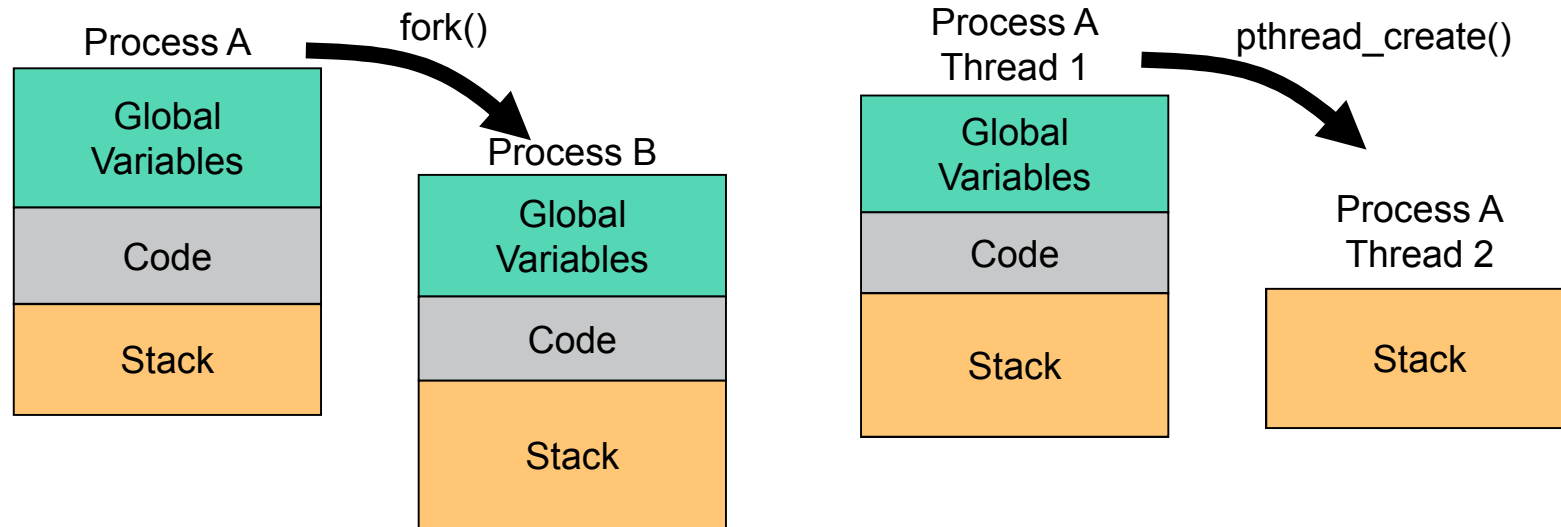
Advantages of thread

- Easy to share information.
 - the memory address space and file descriptors.
- Throughput can be improved.
 - The processing of independent tasks can be interleaved.
- More interactive.
 - The separated threads can deal with user input/output.

Threads concepts

Advantages of thread(cont.)

- The cost for creating a new process is low.



Threads concepts

A thread-specific information

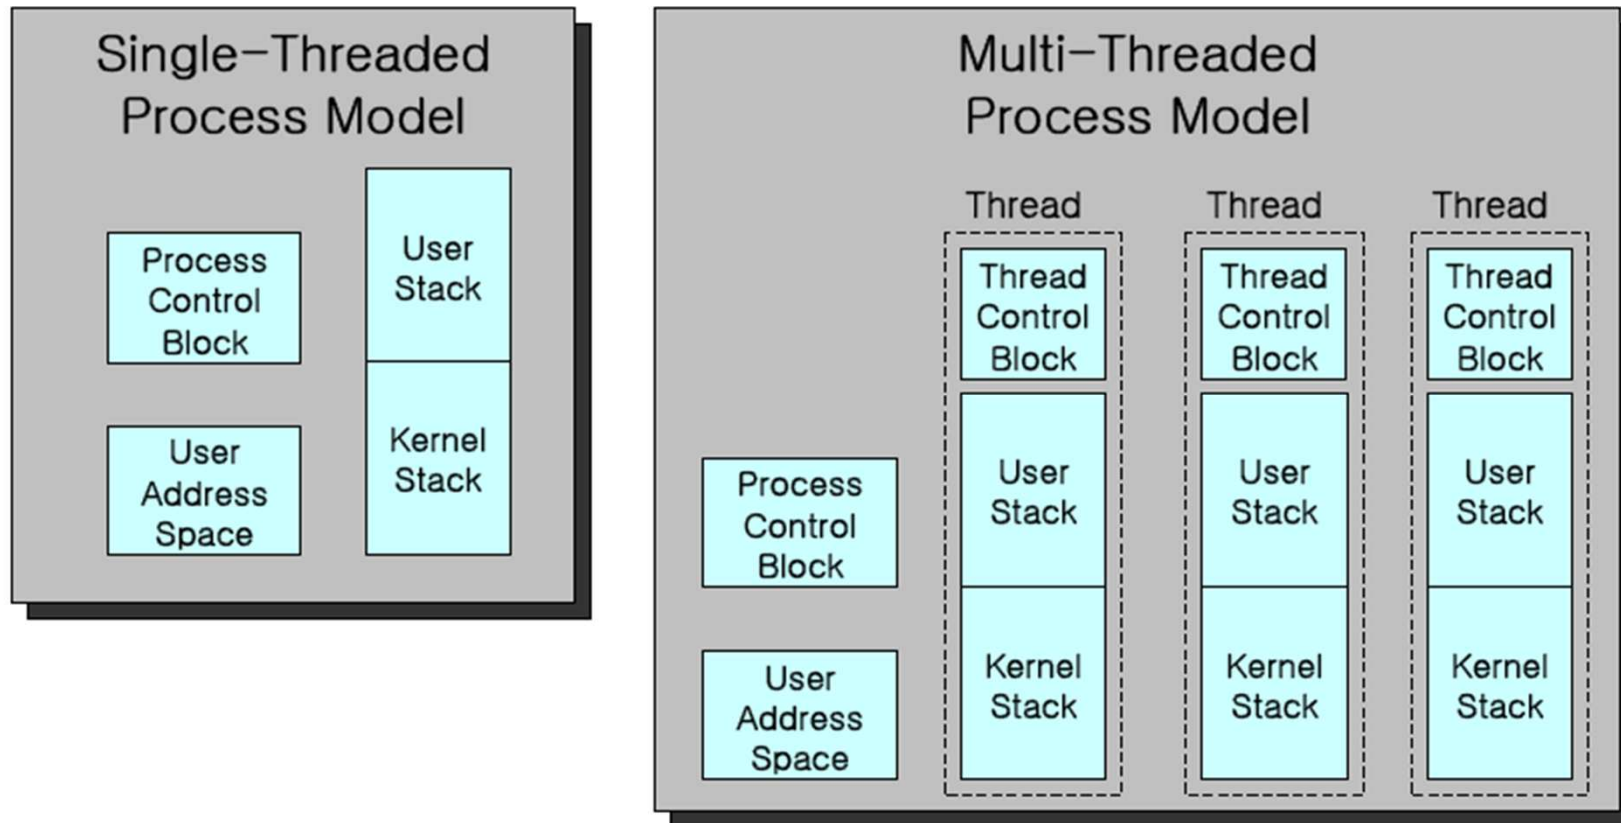
- Thread ID
- Register values
- Stack
- Scheduling priority
- A signal mask

Sharable information among threads in a process

- Text section
- Global data
- Heap
- File descriptor

Threads concepts

❏ Process vs. thread



Posix thread

- ❏ What is pthread?
 - IEEE POSIX 1003.1c standards
- ❏ Pthread naming convention
 - pthread_
- ❏ Compiling pthread program
 - \$ gcc -lpthread xxx.c

Thread identification

Thread ID

- Identifier of thread (similar to process ID.)
- A thread ID is represented by pthread_t data type.
 - Unsigned long integer in Linux.
 - A pointer to the pthread structure in FreeBSD.

Thread identification

```
#include <pthread.h>
```

```
pthread_t pthread_self(void);
```

Returns: the thread ID of the calling thread

 Obtain its own thread ID.

```
#include <pthread.h>
```

```
int pthread_equal(pthread_t tid1, pthread_t tid2);
```

Returns: nonzero if equal, 0 otherwise

 Compare two thread IDs

Thread creation

```
#include <pthread.h>
```

```
int pthread_create(pthread_t *tidp,  
                  const pthread_attr_t *attr,  
                  void *(*start_rtn)(void), void *arg);
```

Returns: 0 if OK, error number on failure

Create a new thread.

- *tidp* is the ID of the newly created thread.
- Execute *start_rtn* with *arg* as its argument.
- *attr* is set to NULL for the default attributes.

Thread creation

Example

```
#include "apue.h"
#include <pthread.h>

pthread_t ntid;

void printids(const char *s)
{
    pid_t    pid;
    pthread_t tid;

    pid = getpid();
    tid = pthread_self();
    printf("%s pid %u tid %u (0x%x)\n", s, (unsigned int)pid,
        (unsigned int)tid, (unsigned int)tid);
}
```

Thread creation

Example(cont.)

```
void *thr_fn(void *arg)
{
    printids("new thread: ");
    return((void *)0);
}

int main(void)
{
    int    err;

    err = pthread_create(&ntid, NULL, thr_fn, NULL);
    if (err != 0)
        err_quit("can't create thread: %s\n", strerror(err));
    printids("main thread:");
    sleep(1);
    exit(0);
}
```

Thread creation

실행

In Solaris

\$./a.out

When a thread is created, there is no guarantee which runs first.

main thread: pid 7225 tid 1 (0x1)

new thread: pid 7225 tid 4 (0x4)

\$

In FreeBSD

\$./a.out

FreeBSD uses a pointer to the thread data structure for its thread ID.

main thread: pid 14954 tid 134529024 (0x804c000)

new thread: pid 14954 tid 134530048 (0x804c400)

\$

In Linux

\$./a.out

Linux does not have a separate system call for thread creation. It uses **clone()**.

new thread: pid 6628 tid 1026 (0x402)

main thread: pid 6626 tid 1024 (0x400)

\$

Thread termination

- ❏ If any thread within a process call `exit()`?
 - The entire process terminates.

- ❏ A single thread can exit without terminating the entire process.
 - The thread **can simply return** from the start routine.
 - The thread can be **canceled by another thread** in the same process.
 - The thread can call **`pthread_exit()`**.

Thread termination

```
#include <pthread.h>

void pthread_exit(void *rval_ptr);
```

 Terminates a calling thread.


- *rval_ptr* is available to other threads in the process calling the `pthread_join()`.

Thread termination

```
#include <pthread.h>
```

```
int pthread_join(pthread_t thread, void **rval_ptr);
```

Returns: 0 if OK, error number on failure

 Suspends execution of the calling thread until the target thread terminates.

- It is similar to **wait()**.
- *rval_ptr* argument
 - If not NULL, it contains the exit status of the target thread.
 - If we're not interested in a return value, it is set to NULL.

Thread termination

Example

```
#include "apue.h"
#include <pthread.h>

void *thr_fn1(void *arg)
{
    printf("thread 1 returning\n");
    return((void *)1);
}

void *thr_fn2(void *arg)
{
    printf("thread 2 exiting\n");
    pthread_exit((void *)2);
}
```

Thread termination

Example(cont.)

```
int main(void)
{
    int      err;
    pthread_t tid1, tid2;
    void      *tret;

    err = pthread_create(&tid1, NULL, thr_fn1, NULL);
    if (err != 0)
        err_quit("can't create thread 1: %s\n", strerror(err));

    err = pthread_create(&tid2, NULL, thr_fn2, NULL);
    if (err != 0)
        err_quit("can't create thread 2: %s\n", strerror(err));
```

Thread termination

Example(cont.)

```
err = pthread_join(tid1, &tret);
if (err != 0)
    err_quit("can't join with thread 1: %s\n", strerror(err));
printf("thread 1 exit code %d\n", (int)tret);

err = pthread_join(tid2, &tret);
if (err != 0)
    err_quit("can't join with thread 2: %s\n", strerror(err));
printf("thread 2 exit code %d\n", (int)tret);

exit(0);
}
```

Thread termination

실행

```
$ ./a.out  
thread 1 returning  
thread 2 exiting  
thread 1 exit code 1  
thread 2 exit code 2  
$
```

Thread termination

```
#include <pthread.h>
```

```
int pthread_cancel(pthread_t tid);
```

Returns: 0 if OK, error number on failure

 Cancel another thread in the same process.

- Cause the thread with *tid* to behave as if it had called `pthread_exit()`.
- It doesn't wait for the thread to terminate; it merely makes the request.

Thread termination


Comparison of process and thread primitives

Process primitive	Thread primitive	Description
fork	pthread_create	Create a new flow of control
exit	pthread_exit	Exit from an existing flow of control
waitpid	pthread_join	Get exit status from flow of control
getpid	pthread_self	Get ID for flow of control
abort	pthread_cancel	Request abnormal termination of flow of control

Thread synchronization

❏ A synchronization example

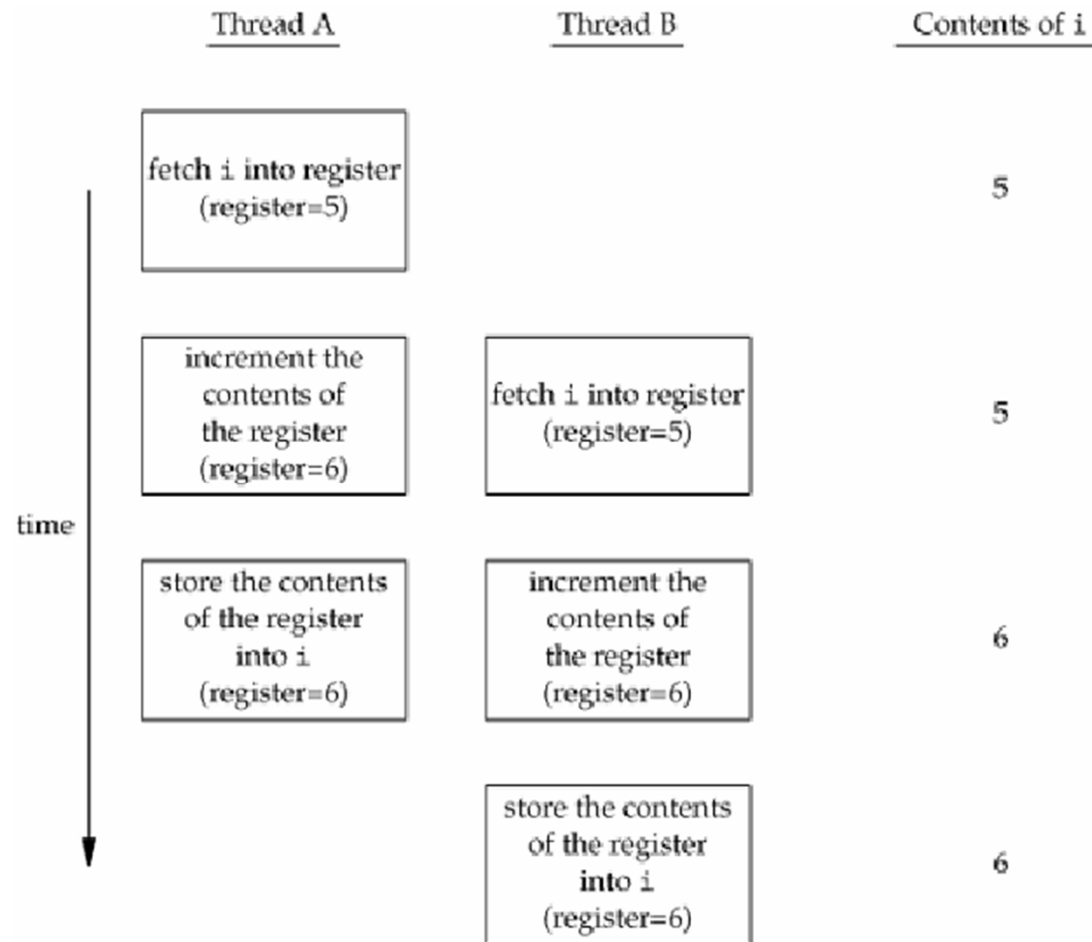
- Two threads try to modify the same variable simultaneously.
 - *i* is initialized to 5.
 - Two threads perform “*i*++;”, respectively.



- Simple statement “*i*++;” consists of several instructions.
 - Read the memory location into a register.
 - Increment the value in the register.
 - Write the new value back to the memory location.

Thread synchronization

A synchronization example



Mutex

Mutex

- It is a lock that we **set (lock)** before accessing a shared resource and **release (unlock)** when we're done.
- While it is set, any other thread that tries to set it will **block until it is released**.
- If more than one thread is blocked when the mutex is unlocked, then **only one** will be able to set the lock.

Mutex

```
#include <pthread.h>
```

```
int pthread_mutex_init(pthread_mutex_t *restrict mutex,  
                        const pthread_mutexattr_t *restrict attr);
```

Return: 0 on success; otherwise, an error number

Initialize a mutex.

- initialize the mutex referenced by *mutex* with attributes specified by *attr*.
- *attr* is generally set to NULL.

Static initialization

- `pthread_mutex_t mutex = PTHREAD_MUTEX_INITIALIZER;`

Mutex

```
#include <pthread.h>
```

```
int pthread_mutex_destroy(pthread_mutex_t *mutex);
```

Return: 0 on success; otherwise, an error number

 Destroy a mutex.

- destroy the mutex object referenced by *mutex*.

Mutex

```
#include <pthread.h>
```

```
int pthread_mutex_lock(pthread_mutex_t *mutex);
```

```
int pthread_mutex_unlock(pthread_mutex_t *mutex);
```

Return: 0 on success; otherwise, an error number

Lock and unlock a mutex.

- If the mutex is already locked, the calling thread shall block until the mutex becomes available.

Mutex

Example (mutex_test.c)

```
#include <pthread.h>
#define NLOOP 50

int counter;                                // This is incremented by two threads
pthread_mutex_t counter_mutex = PTHREAD_MUTEX_INITIALIZER;
void *doit(void *);

int main(int argc, char **argv)
{
    pthread_t tidA, tidB;
    pthread_create(&tidA, NULL, &doit, NULL);
    pthread_create(&tidB, NULL, &doit, NULL);

    // wait for both threads to terminate
    pthread_join(tidA, NULL);
    pthread_join(tidB, NULL);

    exit(0);
}
```

Mutex

Example

```
void *doit(void *vptr)
{
    int i, val;

    for (i = 0; i < NLOOP; i++)
    {
        pthread_mutex_lock(&counter_mutex);
        val = counter;
        printf("%d: %d\n", pthread_self(), val + 1);
        counter = val + 1;
        sleep(1);
        pthread_mutex_unlock(&counter_mutex);
        sleep(1);
    }
    return(NULL);
}
```

Mutex

실행

```
[tskim@oslab test]$ gcc -o mutex_test -lpthread mutex_test.c
[tskim@oslab test]$ ./mutex_test
1093278016: 1
1103767872: 2
1093278016: 3
1103767872: 4
1093278016: 5
1103767872: 6
1093278016: 7
1103767872: 8
1093278016: 9
1103767872: 10
1093278016: 11
...

1093278016: 99
1103767872: 100
[tskim@oslab test]$
```